

Introduction

Representation Learning

Vladimir Zaigrajew

2025-02-26



Vladimir Zaigrajew - vladimir.zaigrajew.dokt@pw.edu.pl

Tymoteusz Kwieciński - tymoteusz.kwiecinski.stud@pw.edu.pl

You can find us in Room 303, MINI, *PW*

Remember every information you can find on our Github Repo:



Figure 1: QR code to course Github Repo



Figure 2: QR code to our Github Repo

Let's chat on slack or discord, because I don't like Teams :/.

In machine learning, we usually want to predict some value $y \in Y$ given some data $x \in X$: we want to learn a function $f : X \rightarrow Y$

| Domain | Task | Example Output |
|------------|--|---|
| Image | Segmentation, Detection, Classification | Class labels, Bounding boxes |
| Text | Sentiment Analysis, Next Word Prediction | Sentiment scores, Text generation |
| Multimodal | Image Description, Image Generation | Generated images, Text descriptions |

Instead of learning a direct mapping $f : X \rightarrow Y$ from input to output, representation learning approach split the problem into two parts: learning a representation $g : X \rightarrow Z$ that transforms raw data into a meaningful feature space, followed by learning a classifier/predictor $h : Z \rightarrow Y$.

X and Y don't have to be from the same domain.



The complete model can be expressed as $f(x) = h(g(x))$

This approach brings several advantages. Most importantly, we can learn representations without labels (unsupervised/self-supervised), reducing the need for manual labeling. With good representations, simpler classifiers can then be used for different tasks, making learning faster and more efficient.



Representation learning transforms complex data into a simpler format that captures important features. Think of face recognition: instead of working with raw pixels, we learn meaningful features like pose and identity, making the recognition task easier.

Learning Representation - Example

Traditional




With Representation Learning



What makes a good representation? I


- **Smoothness:** Similar inputs should have similar representations. If $x_1 \approx x_2$, then $g(x_1) \approx g(x_2)$. This fundamental property ensures that our representations are stable and meaningful.
- **“Less” supervised learning:** Good representations can be learned with minimal supervision, enabling self-supervised and semi-supervised approaches.
- **Invariances/Equivariance/Coherence:** Generally, small temporal/spatial changes should result in similar representations.
Domain specific: image representations should be invariant under transformations like rotations, color jitter etc.



"cat" = "cat"


(a) Invariance¹

- Example: Face recognition should be invariant to lighting changes
- $h(g(x)) = h(g(T(x)))$, where T is some not important transformation



(b) Equivariance¹

- Example: If you rotate an image, the features should rotate similarly
- If T is a transformation, then $g(T(x)) = T(g(x))$



(c) Coherence/Smoothness

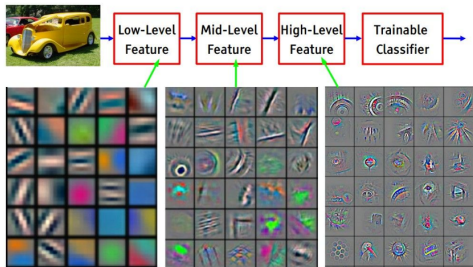
- Close inputs should map to close representations
- Important for generalization and robustness
- If $x_1 \approx x_2$, then $g(x_1) \approx g(x_2)$

¹Source: <https://towardsdatascience.com/sesn-cec766026179/>


- **Smoothness:** Similar inputs should have similar representations. If $x_1 \approx x_2$, then $g(x_1) \approx g(x_2)$. This fundamental property ensures that our representations are stable and meaningful.
- **“Less” supervised learning:** Good representations can be learned with minimal supervision, enabling self-supervised and semi-supervised approaches.
- **Invariances/Equivariance/Coherence:** Generally, small temporal/spatial changes should result in similar representations.
Domain specific: image representations should be invariant under transformations like rotations, color jitter etc.
- **Multiple explanatory factors:** Representations should capture diverse aspects of the data, so that the representation is useful for many different tasks.

What makes a good representation? II

- **Natural clustering:** Representations should reflect natural categories in data, which aligns with human-interpretable groupings.
Example: In vehicle classification, representations should cluster vehicles by type (car, truck, motorcycle) rather than by brand.
- **Hierarchical explanatory factors:** Features organized from concrete to abstract, starting from low-level (edges, colors) to high-level (objects, scenes).
- **Disentangle underlying factors:** Each dimension represents distinct meaningful features, making it easier to understand and manipulate the representation.
- **Sparsity:** For any input x , only few factors are relevant \Rightarrow most dimensions of $g(x)$ should be zero.




(a) Natural clustering and Hierarchical explanatory factors



(b) Disentangle underlying factors and Sparsity


- Principal Component Analysis (PCA)
- Independent Component Analysis (ICA)
- Linear Discriminant Analysis (LDA)
- Multidimensional Scaling (MDS)
- ISOMAP



2

²"Principal component analysis." Wikipedia, Wikimedia Foundation, 10 Apr. 2023, en.wikipedia.org/wiki/Principal_component_analysis.

- Principal Component Analysis (PCA)
- Independent Component Analysis (ICA)
- Linear Discriminant Analysis (LDA)
- Multidimensional Scaling (MDS)
- ISOMAP
- *t-SNE* (t-Distributed Stochastic Neighbor Embedding)
- *UMAP* (Uniform Manifold Approximation and Projection)




3

³McInnes, Leland, John Healy, and James Melville. "UMAP: uniform manifold approximation and projection for dimension reduction. arXiv." arXiv preprint arXiv:1802.03426 10 (2018).


Manifold Learning

The manifold hypothesis states that real-world high-dimensional data tends to lie on or near a lower-dimensional manifold.


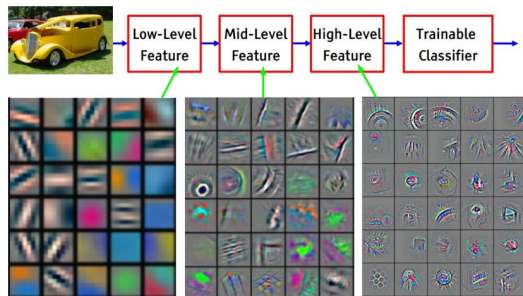


Neural networks learn representations at multiple levels:

- Before Input Layer (pixels, words, etc.)
- After Input Layer
- Hidden Layer
 - Progressively more abstract features
 - Combine and transform earlier representations
 - Different layers capture different aspects
- Final Layers (task-specific representations)

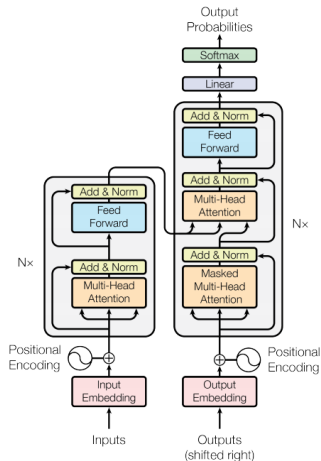


CNNs demonstrate hierarchical representation learning



```
Resnet18 layer layer1 representation shape torch.Size([1, 64, 56, 56])
Resnet18 layer layer2 representation shape torch.Size([1, 128, 28, 28])
Resnet18 layer avgpool representation shape torch.Size([1, 512, 1, 1])
```

Transformers Representations



```
import torch
import transformer_lens
```

```
# Load a model (eg GPT-2 Small)
model = transformer_lens.HookedTransformer.from_pretrained(
    "gpt2-small")
```

```
# Run the model and get logits and activations
input_str = "Tancowala ryba z rakiem, ryba z rakiem"
with torch.no_grad():
    logits, activations = model.run_with_cache(input_str)
block_to_visualize = activations['blocks.11.ln1.hook_normalized']
Loaded pretrained model gpt2-small into HookedTransformer
```

```
Input text: Tancowala ryba z rakiem, ryba z rakiem
Logits shape: torch.Size([1, 20, 50257])
After embedder layer: torch.Size([1, 20, 768])
After Layer 11: torch.Size([1, 20, 768])
```

4

⁴Vaswani, Ashish, et al. "Attention is all you need." Advances in neural information processing systems 30 (2017).

Key takeaways about representation learning:

- ① Data representation matters
 - Raw data (images, text) is often not optimal for ML models
 - Good representations make learning easier
- ② Multiple levels of abstraction
 - From raw features to high-level concepts
 - Different representations serve different purposes
- ③ Future directions
 - Self-supervised learning (**Next week**)
 - Multi-modal representations
 - More interpretable representations